

Assignment

Database Management Systems

Instructor: Nitin Awathare

Overview

This assignment requires you to implement the stated objective on the **toyDB** database, for which the necessary files have been provided. The toyDB system is a simplified database that contains the two fundamental lower layers of a database management system (DBMS).

The first is the **Paged File (PF) layer**, which is responsible for handling storage at the level of individual pages within a file. This layer abstracts away low-level file operations and provides an interface for managing fixed-size data pages efficiently.

The second is the **Access Method (AM) layer**, which builds upon the PF layer to support B⁺ tree-based access structures. This layer provides efficient indexing and retrieval operations, enabling structured access to data.

Detailed descriptions of the interfaces exposed by these two layers are available in the provided documentation files (**pf.ps** and **am.ps**). In addition, the C source code for both the PF and AM layers has been included, giving you the necessary foundation to understand their implementation and extend the functionality as required for this assignment.

In addition to the database code and documentation, the assignment package also includes test data that students can use to carry out and validate their implementations. This data is bundled in the file **data.tar.gz**, which, once extracted, provides a collection of databases specifically designed for use with the assignment.

The datasets are stored as simple text files, each representing a table of information. These files contain dummy data that has been carefully prepared solely for the purpose of this assignment. The structure of the data closely resembles realistic academic information, thereby allowing students to practice with data that feels practical while still being simplified for instructional use.

For instance, one of the files, the *courses* file, contains records of all courses offered across different departments. Each record includes fields such as the course number, the course name, and whether the course is intended for postgraduate (PG) or undergraduate (UG) students. The records are stored in a plain-text format, with individual fields separated by semicolons (;).

Objectives

1. Page Buffering for the PF Layer

Implement page buffering for the PF layer with support for two page replacement strategies — **LRU** and **MRU** (the latter being particularly useful for sequential file access). The replacement strategy should be selectable when opening a file. Each page must maintain a *dirty flag* to track modifications, with an explicit call provided to mark a page as updated. The buffer pool size should be configurable as a parameter.

Additionally, the system must collect and report statistics, including:

- Number of pages accessed (read/write)
- Logical and physical I/O counts for different mixtures of read and write queries

You must plot a graph where the X-axis represents different mixtures of read/write queries and the Y-axis represents the collected statistics.

2. Slotted-Page Structure

Build a **slotted-page structure** on top of the PF layer to store variable-length records, using it to manage student data. The system should support:

- Record insertion
- Record deletion
- Sequential scanning

Additionally, gather performance metrics such as *space utilization* and compare it with the static management of records. Your results should be presented in a table, taking into account different maximum record lengths during static record management.

3. Index Construction via the AM Layer

The AM layer can be used to build an index on a file. This can be done in two ways:

- a. Starting with a file that already contains records and creating the index in a single operation.
- b. Beginning with an empty file and building the index incrementally through multiple inserts.

Evaluate both approaches for constructing an index on the *Student* file using the *roll-no* field as the key. If the file is pre-sorted, more efficient bulk-loading techniques for index construction are available (as discussed in class) — implement one such method and compare its performance (*query completion time*, *number of pages accessed*) with the two indexing approaches listed above.

Tips for Testing

Before beginning the implementation, you are advised to familiarise yourself with the code hierarchy of the ToyDB database and test it using basic read queries without making any modifications.

Submission Instructions

1. The code should follow programming etiquette with **appropriate comments**.
2. Add a **README** file which includes a description of the code and gives detailed steps to compile and run the code.